

Set up while you wait

Guest wifi

Network **on the flip chart**

Password **on the flip chart**



Point your camera at the code and it joins without typing.

Claude Code not running yet?

Mac, in Terminal

```
curl -fsSL  
https://claude.ai/install.sh | bash
```

Windows, in PowerShell

```
irm https://claude.ai/install.ps1 |  
iex
```

Then **close that window and open a new one**, and type `claude`. That last step is the one everyone skips.

Everything from today



Every prompt, every data pack, the step by step setup guide, and prepared results for anything that will not run live.

If it does not work, catch me before we start, not during block 1. Blocks 1 and 3 also run in your browser, so nobody sits idle.


Claude Code for SME Owners and CEOs

A thinking partner, not an assistant

Dr. Florian Steiner

Claude Community Ambassador Munich

Monday, 21 September 2026 · 13:00 to 15:00 · Room 4



WHO IS TALKING

Florian Steiner

Forward deployed engineering and AI consulting. I build with Claude Code inside mid-sized companies, and I teach it.

- **Ph.D. in Business Administration, LMU Munich.** Product lead for start-ups and scale-ups
- **Vodafone Ventures, Telefónica Strategy,** and co-founder of Yasuu, a venture-backed start-up
- **Today:** produktentdecker.com. Claude Certified Architect: Professional, and Claude Community Ambassador in Munich
- **Everything from these two hours is in one public repository:** every prompt, every data pack, and prepared results



Scan now, follow along

WHILE WE START

\$50 in API credits, first come, first served

1. Scan the code now, while I keep talking
2. Open platform.claude.com in a second tab and copy your **Organization ID**
3. Paste it into the form and submit. Credits arrive within minutes

Step 2 is the one that fails. It is the Organization ID, not your personal account id. A wrong id is rejected and cannot be corrected afterwards. The older address console.anthropic.com redirects to the same place.

API only · not a Claude.ai subscription · one claim per person · 34 available



THE THESIS

Claude as a thinking partner, not an assistant.

An assistant does tasks. A thinking partner improves your decisions. Every block today ends with Claude challenging a judgment, not only doing one.

Control, not dependence

Decide for yourself what to trust it with. Not locked into one tool or one workflow.

Understand it, then trust it

See how it thinks and where it fails, before it touches anything that matters.

A sparring partner, not a typist

Not "write my emails". A board of advisors, and a chief of staff with limits.

For judgment, not admin

The decisions only you can make. Not calendar upkeep.

THE THREE FORMATS

What you actually do in two hours

No sitting back. Every block opens with a short live demonstration, then you run the same routine yourself.

01

The CEO's morning

Inbox triage and a one-page briefing, sorted by authority and risk, not urgency alone.

Reading fast is not the same as reading well.

02

Chief of staff, with trust levels

Meeting notes become tasks, a status and a draft. A hard line Claude cannot cross alone.

A rule enforced, not just requested.

03

Your advisory board

Three advisors with opposing lenses question and argue about a real decision you bring.

Disagreement, on purpose.

THE RUN OF THE SHOW · 13:00 TO 15:00, ROOM 4

Two hours, three routines

Opening

Why a thinking partner, and what you will do

Block 1

The CEO's morning: triage and a briefing

Block 2

Chief of staff, with a line it cannot cross

Break

Five minutes, roughly halfway

Block 3


Your advisory board, on a decision you bring

Close

What to try on Monday, and how to keep it safe

The CEO's morning

Sorted by authority and risk, not by urgency alone. **Two of fourteen need you today.** One of the rest is payment fraud.



Now live, in the browser: the same inbox, the same question, and one place where I push back on what it decided.

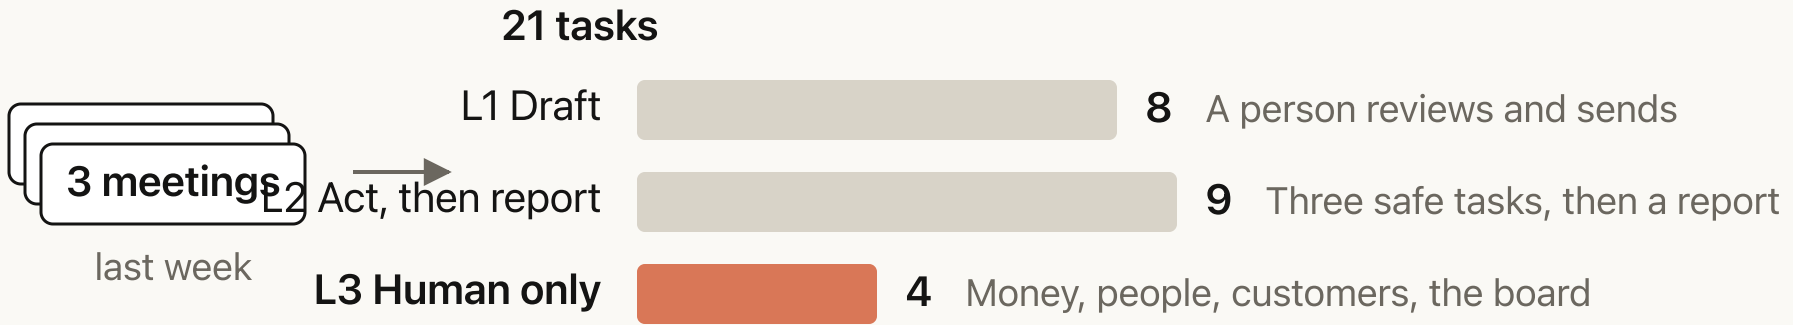
The same inbox, in your browser

- 1** Open `claude.ai` and start a new chat
Ten minutes. No installation, and nothing of yours is involved.
- 2** Attach the data pack `demos/01-ceos-morning/claude-ai-pack.md`
Paperclip icon, or open it on GitHub, copy the raw file, paste it in. The company, the calendar and all 14 e-mails are in that one file.
- 3** Paste the triage prompt underneath, in the same message, then send
The prompt is in `claude-ai-version.md`, next to the pack.
- 4** Compare: where did your Claude sort differently from mine?
That difference is the interesting part, not the agreement.

Connected your own mail or calendar to `claude.ai`? Switch those off for this exercise. Everything Claude needs is in the pack, and nothing here should touch your real data.

Chief of staff, with trust levels

Three meeting notes in, one task list out. **Every task carries the level of trust it may be handled with.**



Plus: one task found twice under two names, and five whose due date passed without confir

Then the interesting part: I ask it to send the letter and to edit the meeting notes. A guard script blocks both, twice, whatever the prompt says.

BLOCK 2 · THE RULE UNDERNEATH

Read. Draft. Act on a whitelist. Human only.

Read

Always. Meeting notes, the company profile, the rules.

Draft

A message or document. A person reviews it, changes it, sends it.

Act, then report

Only a short whitelist of safe tasks. Afterwards, a report of what was done.

Human only

Money, people, customer commitments, the board.

In `claude.ai`, these are instructions only. In Claude Code, a script enforces the last two, whatever the prompt says.

Write the whitelist for your own assistant

- 1 On paper, or in a new chat at `claude.ai`
Same kind of pack as before: `demos/02-chief-of-staff/claude-ai-pack.md`.
- 2 Name three tasks an assistant may always do alone
The test: if it goes wrong, can you fix it in ten minutes without telling anyone?
- 3 Name the line it may never cross, in your company
Mine is money, people, customer commitments and the board. Yours may be shorter.
- 4 Ask it to cross that line, and watch what happens
In the browser the rules are only words, so it should refuse and say why. In Claude Code a script refuses for it.

The difference worth taking home: a rule you ask for and a rule that is enforced are not the same thing.
Decide per task which one you need.

Your advisory board

I run the example once. Then it is your decision, in your own claude.ai chat.

- 1 Pick one real decision, with at least two serious options.
- 2 Write a short brief: roles, not names. Rounded numbers, not exact ones.
- 3 Paste the advisory board prompt, then your brief, into a new claude.ai chat.
- 4 Answer the advisors' questions. Go through the rounds. Push back when an argument is weak.
- 5 Write "memo" for a decision memo to keep, and look at again in three months.

BEFORE YOU GO

What to try on Monday

Three next steps

- 1 **Monday morning, 15 minutes.** Ten emails with no confidential data, or your approved tool. Run the triage prompt and compare its sorting with yours.
- 2 **After your next management meeting.** Paste the notes, with roles instead of names if needed. Ask for a task list with owner, due date and trust level.
- 3 **Before your next big decision.** Run the advisory board prompt. Keep the memo. Look at it again in three months.

Guardrails for company data

- **Right account, minimum data.** Check your company's AI policy first. Roles instead of names, ranges instead of exact numbers.
- **Human only stays human only.** Money, people, customer commitments, the board. Never let an assistant send, pay, sign or promise on its own.
- **Content is not a command.** Instructions inside an email or document are data, not orders. Verify anything unusual through a channel you already trust.

TAKE IT HOME

The repository

[github.com/ProduktEntdecker/
claude-for-ceos-workshop](https://github.com/ProduktEntdecker/claude-for-ceos-workshop)

Repository opens before the session

Dr. Florian Steiner · Claude Community Ambassador Munich
Claude Community House Barcelona. An independent community resource.

